

Knowledge Representation

by Adia Khalid

Summer semester 2026

Lecture notes by Liliana Sanfilippo

April 2026

Contents

I	Foundation of Knowledge Representation	5
1	Introduction to KR: Definition, Scope, Motivation (Use Cases), and History	7
1.1	Introduction	7
1.2	Taxonomy of Representation Schemes	9
1.3	Distinguishing Features of KR Systems	12
1.4	Applications of Knowledge Representation	15
1.5	Future Directions	17
1.6	Conclusion <i>brought to you by claude.ai</i>	20
1.7	Course structure	21
1.8	Formalities	21
2	Logical Foundations: First-Order Logic (FOL)	23
2.1	Syntax	23
2.2	Semantics	24
2.3	Inference	24
2.4	Proof Systems & Completeness	25
3	Logic Programming and Automated Reasoning	27
4	Structured Knowledge Representation	29
II	Knowledge Graphs	31
5	Start into semantic networks and knowledge graphs reasoning	33
6	The Resource Description Framework (RDF)	35
7	Knowledge Graph Reasoning	37
8	Knowledge Acquisition from Data	39
9	Representation Learning	41
III	Knowledge Representation in Machine Learning	43
10	Knowledge Graph Embeddings	45
11	Graph Neural Networks for Knowledge Representation	47

IV Applications	49
12 Knowledge Tracing	51
13 Knowledge in Large Language Models	53
14 KG- Aware Applications	55
Backmatter	i
Figures	i
Index	iii
Definitions	iii
References	vii

Part I

Foundation of Knowledge Representation

CHAPTER 1

Introduction to KR: Definition, Scope, Motivation (Use Cases), and History

1.1 | Introduction

What is Knowledge Representation?

Definition 1.1 Knowledge Representation

KR is a central problem in Artificial Intelligence that focuses on developing sufficiently precise notations for representing knowledge

The core problem with KR is that the real world with its complexity, ambiguity and dynamism can never be fully represented in a computer-based system. Nevertheless, intelligent systems must be able to function effectively using partial, incomplete and sometimes contradictory information. [4]

Why is Knowledge Representation needed?

The problem of data meaninglessness

To understand the need for KR, let's consider a practical scenario:

Example 1.1 A student database.

A hypothetical database contains the following entries:

Name	Major	CGPA	Research	...
Alice Smith	CS	3.8	Genomics	...
Bob John	Biology	3.8	Oncology	...
Carol Li	Zoology	3.7	Animal colonisation	...

Without knowledge representation, this data is disjointed. The entries exist in isolation across thousands of files. A database query system can only perform exact keyword matches:

- Query 1: "Show me all AI researchers" → No results (not explicitly mentioned in the data)
- Query 2: "Which students could work on this patent?" → No answer possible
- Query 3: "What are the emerging trends?" → No analysis possible

Data is meaningless without context and relationships. It does not allow for intelligent reasoning or deeper analysis.

What knowledge representation makes possible

With knowledge representation, we can structure the data and transfer isolated records into interconnected, queryable knowledge using:

- **ontologies**: Explicitly specify which concepts exist (student, researcher, field of research) and how they are categorised.
- Knowledge graphs: Link entities together: Alice → researcher → AI

This allows for semantic searches, automatic relationship discovery and reasoning across all connection.

Example 1.2 • Semantic search

- “Find students researching machine learning?”
- “Show me all AI researchers”: The system can use transitive relationships to find relevant students.
- “Which students are a good fit for this patent?": The system can match skills and areas of interest.
- Automatic relationship discovery
 - “What are emerging research trends?”
 - Recognising patterns across all data
 - Identifying emerging trends through interconnected information

Data	Information	Knowledge
Raw, unprocessed facts	Structured and contextualised data	Information combined with experience and reasoning rules
e.g. CGPA: 3.8	e.g. Alice has a CGPA of 3.8 in Computer Science	e.g. Alice could work on this AI patent because she has a strong foundation in computer science and is interested in genomics

Table 1.1: Data vs. Information vs. Knowledge

KR enables machines to make informed decisions with fewer rules.

KR Perspective

Traditional ML systems follow a black-box approach which means the decision-making process is not transparent and the reasoning behind predictions is not accessible.

- Generalisation: Requires continuous retraining with new data
- Explainability: Difficult or impossible to explain decisions
- Data flexibility: Highly dependent on large amounts of training data

Knowledge-Based AI utilises structured, **explicit knowledge**.

Definition 1.2 explicit knowledge

Knowledge that is formally encoded in symbols, rules or graphs and can be processed by machines

- Generalization: explicit knowledge transfers to new problems without retraining from scratch[3]
- Explainability: encoded knowledge enables transparent reasoning and human-understandable decisions[1]
- Data flexibility: Modifiable

Representation Schemes

Definition 1.3 representation scheme

A representation scheme is a formal notation used to define a knowledge base (Mylopoulos 1980). A knowledge base contains facts and rules and serves as a model of the real world. A representation scheme is therefore the ‘language’ in which we express knowledge, much like programming languages define the syntax and semantics in which we write instructions.

Example 1.3 Pac-Man ghost AI has a knowledge-based approach: We code ghost behaviour as explicit rules.

1. IF player_nearby THEN chase
2. IF frightened THEN flee
3. ELSE patrol

Decisions are transparent, explainable, and modifiable without recomputing. (Ogland 2016)

Challenges in Knowledge Representation

Knowledge representation involves designing formal structures that allow computers to store, interpret, and reason about information and every design decision has advantages and caveats.

Choosing the representation is one of the design challenges. **Semantic networks** are intuitive and visually understandable but have limited inference. **First-order logic** on the other hand is expressive, but has computationally costly reasoning.

Desired properties can conflict. Systems should ideally be expressive enough to capture rich real-world knowledge while remaining efficient enough to reason in practical time. Tradeoffs are between expressiveness and tractability, and Human readability and machine efficiency.

Application-driven constraints arise from, different domains requiring different forms of representation. E.g. In natural language processing (NLP), systems must handle ambiguity and context. In medical diagnosis, representations must manage uncertainty, causal relationships, and strict correctness requirements. In question answering, systems need fast retrieval combined with reasoning over large amounts of structured and unstructured data

The core idea of Knowledge Representation is not just a data structure. It plays multiple roles simultaneously and those roles can conflict with each other:

- Surrogate for the world: Acts as a substitute for real world objects and relationships
- Ontological commitments: Defines what exists and how things are categorized
- Theory of reasoning: Provides foundation for logical inference and deduction
- Medium for efficient computation: Enables performant computational processing
- Medium of Human expression: Allows humans to express knowledge in understandable ways

1.2 | Taxonomy of Representation Schemes

Representational schemata can be divided into four basic categories:

- **Declarative representations**

- Logical representation schemes
- Semantic (network) schemes
- Procedural representations
- Hybrid approaches
- Frame-based representations

Declarative representations

Definition 1.4 declarative representation

Declarative representations describe facts about the world as logical statements

Declarative representations focus on truth and semantics and enable deductive reasoning. The schema is independent of how the knowledge is used.

There are two subtypes:

1. Logical representation schemes

Definition 1.5 logical representation scheme

Representation schema based on formal logic and represents knowledge through logical formulas. It is strongly based on semantics and inference

Example 1.4 First order Logic



Figure 1.1: This is Tweety. What can we infer about Tweety with the information given below? (image from <https://www.pexels.com/photo/pink-cockatoo-perched-in-natural-habitat-29732587/>)

We know that all bird can fly if they are not flightless birds, here generalized as penguins. We know that Tweety is a bird and not a penguin.

More formal:

Birds can fly: $\forall x (Bird(x) \wedge \neg Penguin(x) \rightarrow CanFly(x))$

Facts: $Bird(Tweety)$ and $\neg Penguin(Tweety)$

Therefore we can infer:

$CanFly(Tweety)$

2. Semantic (network) schemes

Definition 1.6 semantic representation scheme

Graph-based representations in which nodes represent entities and edges represent relationships between them. It supports inheritance and classification

Example 1.5




Figure 1.2: Semantic Network

Procedural representations

Definition 1.7 procedural representation

Representation schema that encodes knowledge as procedures or programmes that describe how to use that knowledge. It is efficient but hard to interpret

Example 1.6



Figure 1.3: This is Dr. Dokter. She wants to diagnose her patient. (image from <https://www.pexels.com/photo/female-doctor-in-scrub-suit-7659874/>)

The rules are:

IF (patient has fever) **AND** (patient has cough) **AND** (patient has difficulty breathing) **THEN diagnose:**
possible pneumonia, **recommend** chest x-ray

This allows Dr. Dokter to properly diagnose her patients instead of furthering the antibiotic resistant strand of bacteria!

Hybrid approaches combine declarative and procedural aspects.

Frame-Based Representation

Definition 1.8 frame

Concept designed to provide a structured way to capture the essential aspects of a situation, facilitating easier retrieval and manipulation of information. They are akin to schemas or blueprints that organize knowledge into manageable chunks. Attributes of the represented concept are represented by slots which have facets and default values

Example 1.7 Back to Tweety from fig. 1.1.

```

Frame: Bird
  is_a:      Animal
  has_parts: [wings, beak, feathers]
  can:       fly
  default_color: pink
  locomotion: flying

Frame: Penguin
  is_a:      Bird
  can:       swim
  cannot:    fly (overrides parent default)

```

1.3 | Distinguishing Features of KR Systems

Unlike databases (retrieval) or programs (execution), knowledge bases reuse the same facts in multiple ways:

Inference (rule-based reasoning):

Facts are used to derive new facts by applying rules.

Example 1.8 We remember Dr. Dokter from fig. 1.3.

$$\text{IF fever} \wedge \text{cough} \rightarrow \text{possible infection}$$

$\Rightarrow$ Diagnostic conclusion

It has multiple subtypes:

Definition 1.9 logical inference

Inference that uses formal deduction from known facts

Example 1.9

$$\begin{aligned} \text{IF fever} &\rightarrow \text{infection} \\ \text{fever} &\Rightarrow \text{infection} \end{aligned}$$

Definition 1.10 specialized procedures

Fixed inference patterns, which are often implemented as routines

Example 1.10 Ward A $\rightarrow$ Ward B $\rightarrow$ ICU $\Rightarrow$ Shortest path

Definition 1.11 preferred inference

Reasoning based on typical assumptions, subject to certain reservations

Example 1.11

Birds fly
 Tweety is a bird $\Rightarrow$ Tweety flies
 (except penguins)

Definition 1.12 inductive inference

Generalising examples into patterns or rules

Example 1.12

Data $\Rightarrow$ Pattern
 Symptoms $\Rightarrow$ Diagnosis model

Definition 1.13 abductive inference

Reasoning to arrive at the best explanation. From observations to hypothetical causes

Example 1.13

Observed: fever + cough
 $\Rightarrow$ Best explanation: pneumonia

Access (Direct fact retrieval via queries):

Facts are retrieved directly, much like in a database.

Example 1.14 Query: “Which patients have fever?” $\Rightarrow$ Retrieve patient X

Matching (Similarity-based reasoning):

Facts are used to identify similar patterns and perform case-based reasoning.

Example 1.15

$X \sim \text{Patient Y} \Rightarrow \text{Case-based reasoning}$

Design implications

When designing a system, we must choose which inference mechanisms to support (e.g., shortest-path search vs. rule-based reasoning). Of course, there are trade-offs between expressiveness and efficiency. And different applications require different inference types, e.g. the Pac-Man game (example 1.3) needs heuristic inference for real-time decisions goals.

Observation 1.1 Knowledge bases cannot completely describe the world they model (except artificial microworlds)

Consequences exist for operations, reasoning and design methodology.

Consequences for Operations

Implication 1.1 Regarding observation 1.1

There are different answers to existence queries depending on assumptions.

We can assume a closed world, where everything not explicitly stated in the knowledge base is assumed FALSE or an open world, where we simply cannot conclude anything if it is not explicitly stated in the knowledge base.

Example 1.16 To definition ?? - Closed World Assumption

Knowledge Base: Alice (New York), Bob (London), Carol (Paris)

Query: “Is there a person in Toronto?”

CWA Answer: NO (because Toronto is not mentioned in the KB)

Reasoning: If someone lived in Toronto, the KB would state it.

Example 1.17 To definition ?? - Open World Assumption

Knowledge Base: Alice (New York), Bob (London), Carol (Paris)

Query: “Is there a person in Toronto?”

OWA Answer: Unknown (because Toronto is not mentioned, but also not excluded)

Reasoning: The KB does not mention Toronto, but that does not mean no one lives there.

The choice between CWA and OWA fundamentally influences system behavior and must be made during the design phase.

Consequences for Reasoning

Incompleteness requires special reasoning strategies:

- Assumption of Defaults: In the absence of information, make reasonable assumptions
- Probabilistic Reasoning: Explicitly model uncertainties
- Constraint Propagation: Use constraints to reduce search spaces

Consequences for Design Methodology

Unlike machine learning, where knowledge is implicit in model parameters and retraining is required, in KGs knowledge is explicit and can be extended, revised, and corrected without retraining the entire model.

This distinction is important for practical systems where expert knowledge needs to be continuously updated without expensive computational overhead.

Programming	Knowledge Representation	Machine Learning
“Once and for all” complete specification. Code is written and then “frozen”	Continuous, incremental development. The knowledge base is iteratively extended, revised, and refined	Knowledge is implicitly encoded in model parameters and requires retraining when incomplete

Table 1.2: Differences in knowledge/information acquisition and handling.

Definition 1.14 Knowledge Acquisition Processes

The process of incorporating knowledge into the knowledge base through strategies such as Interactive sessions with experts or Automatic generation from experience

1.4 | Applications of Knowledge Representation

“The KR discipline has developed in waves, with each era developing different approaches to different problems.”

– *claude.ai*

Definition 1.15 **test**
todo

Example 1.18 To definition 1.15 - MYCIN's Approach

Domain: Infectious blood diseases and bacterial infections

Knowledge Base Structure:

- Approximately 600 production rules
- Rules encoded expert medical knowledge
- Format: IF [conditions] THEN [conclusions/actions]

Definition 1.16 **Natural Language Processing**
todo

Natural Language Processing includes Semantic understanding, Machine translation, Question answering and Dialogue systems.

Example 1.19 For **Natural Language Processing**:

- Semantic understanding - “Paris is the capital of France” $\Rightarrow$ understands (`capital_of(Paris, France)`)
- Machine translation - “The doctor treated the patient” $\Rightarrow$ preserves roles across languages (doctor = agent, patient = object)
- Question answering - Q: “Who is the capital of France?” $\Rightarrow$ Query knowledge graph Answer: Paris
- Dialogue systems - User: “I have a fever.” System: “Do you also have a cough?” $\Rightarrow$ Uses context and stored knowledge

Today (in the modern era), we use **knowledge graphs**.

Definition 1.17 **knowledge graph**

Graph-based knowledge representation where the entities are nodes and the relations are edges. Facts (subject, a predicate, and an object) are represented as triples.

Example 1.20 To 1.17

- Google Knowledge Graph (search enhancement) Example: Search “Albert Einstein” $\rightarrow$ panel shows: born in Ulm, Nobel Prize, physicist, spouse Mileva Maric
- Knowledge Vault (knowledge fusion) Example: Automatically merges facts from Wikipedia, Freebase, and news articles into unified knowledge base
- Enterprise Knowledge Management Example: Company stores employee skills, projects, teams $\rightarrow$ helps find experts for new projects

- Recommendation Systems Example: Amazon recommends “Laptop Bag” because you bought “Laptop” → relationship stored in knowledge graph

Knowledge graphs are used for better information retrieval. E.g. for **Information Retrieval (Semantic Search)**s.

Definition 1.18 Information Retrieval (Semantic Search)

todo

Example 1.21 Query: “treatment for viral pneumonia”

Without KG: Documents containing “treatment” + “viral” + “pneumonia”

With KG:

pneumonia $\text{--}[\text{caused_by}]\text{--}$ virus virus $\text{--}[\text{type}]\text{--}$ COVID-19, Influenza, etc. COVID-19 $\text{--}[\text{treated_by}]\text{--}$ antiviral drugs, ventilation

Retrieval: Documents about COVID-19 treatment, antiviral therapies (not necessarily containing the word “pneumonia”)

Knowledge representation techniques are also used in database design for semantic data models, query optimization, constraint management and data integration.

Example 1.22 Semantic Data Models: using an ER diagram to represent Students, Courses, and Enrollments. *Shoutout to claude*

Entities:

- Student (ID, Name, Major, CGPA)
- Course (Course_ID, Name, Credits)
- Enrollment (Student_ID, Course_ID, Grade)

Relationships:

- Student enrolls_in Course
- Course belongs_to Department

Example 1.23 Query Optimization - creating an index on student_id to speed up SELECT queries. *Shoutout to claude*

Query: SELECT student_name FROM Student WHERE enrolled_in_course(CS101)

Plan A (inefficient):

1. Load all students from disk
2. For each student, check all enrollments
3. Filter for CS101

Plan B (efficient):

1. Use index on (Course_ID)
2. Find all enrollments where Course_ID = CS101
3. Load corresponding student records

Example 1.24 Constraint Management - enforcing primary key and foreign key constraints to maintain data integrity. *Shoutout to claude*

Primary Key Constraint: Each student has a unique ID

Foreign Key Constraint: Each enrollment must reference a valid student_ID

Domain Constraint: CGPA must be between 0.0 and 4.0

Business Rule: A student cannot take more than 5 courses per semester

Further, **Knowledge Representation** can be used for program specifications and system interfaces.

Example 1.25 Space Tutor, Folie 37

1.5 | Future Directions

Multi-hop Reasoning (Complex Reasoning)

Many questions require multiple reasoning steps across relationships.

Example 1.26 Query: “What medicine treats diseases caused by virus X?”

Requires:

medicine $\rightarrow$ [treats] $\rightarrow$ disease $\rightarrow$ [caused_by] $\rightarrow$ virus

Path:

Starting from virus_X:

virus_X $\rightarrow$ [causes] $\rightarrow$ disease_Y

disease_Y $\rightarrow$ [treated_by] $\rightarrow$ medicine_Z

Result: medicine_Z

Simple embedding methods struggle with this multi-step logical chain, so complex multi-hop relations require dedicated design.

Promising Directions:

- **Recurrent Relational Path Encoding:**

- Methods like DeepPath use LSTM networks
- The LSTM “encodes” a path like: Person $\rightarrow$ LivesIn $\rightarrow$ City $\rightarrow$ LocatedIn $\rightarrow$ Country
- This encoded path representation can then be used for reasoning

- **GNN-based Message Passing:**

- Graph Neural Networks (GNNs) aggregate information from neighbors
- Multi-hop reasoning through iterative message passing
- Each “hop” collects information from the next node

- **Reinforcement Learning-based Path Finding:**

- An agent learns to navigate through the knowledge graph
- Example agent path: “Marie Curie $\rightarrow$ BornIn $\rightarrow$ Ulm $\rightarrow$... $\rightarrow$ Paris”
- The agent receives rewards when reaching the correct target

Definition 1.19 one-hop

todo

Definition 1.20 **multi-hop**
 todo

Example 1.27 To definitions 1.19 and 1.20 - Simple vs. Multi-hop Reasoning
 A simple 1-hop query can be expressed as:

$$\text{Einstein} \rightarrow \text{worked_at} \rightarrow \text{Princeton}$$

A more complex multi-hop query is: *Where did Einstein's doctoral advisor teach?*
 This requires chaining relations:

$$\text{Einstein} \rightarrow \text{advised_by} \rightarrow \text{Kleiner} \rightarrow \text{worked_at} \rightarrow \text{University of Zurich}$$

This corresponds to a 2-hop reasoning process.

Structured knowledge provides commonsense observations and acts as relational inductive biases (assumptions about the world). Meaning they provide such biases like: “Some relations are more likely than others”. It also boosts commonsense knowledge fusion between symbolic and semantic space. E.g.:

Example 1.28 **Query:** *Can a fish play piano?*

Without Knowledge: A machine learning model could predict this based on training data, but provide no reasoning.

With Knowledge:

$$\begin{aligned} \text{Fish} &\rightarrow \text{hasBodyPart} \rightarrow \text{Fins} \quad (\text{not Hands}) \\ \text{PlayingPiano} &\rightarrow \text{requires} \rightarrow \text{Hands} \end{aligned}$$

Symbolic Rule: If $X \rightarrow \text{hasBodyPart} \rightarrow p$ and $\text{action} \rightarrow \text{requires} \rightarrow p$, then $X \rightarrow \text{can} \rightarrow \text{action}$.

Application:

$$\begin{aligned} \text{Fish} &\rightarrow \text{hasBodyPart} \rightarrow \text{Fins} \\ \text{PlayingPiano} &\rightarrow \text{requires} \rightarrow \text{Hands} \\ \text{Fins} &\neq \text{Hands} \end{aligned}$$

Conclusion:

$$\text{Fish} \rightarrow \text{cannot} \rightarrow \text{PlayPiano}$$

Answer with Reasoning: No, a fish cannot play piano because it lacks hands.

Commonsense knowledge provides strong priors, enabling systems to function with less training data.

Unified Framework

Many KR embedding methods are mathematically equivalent under transformation. Actually, several representation learning models are verified as equivalent.

Example 1.29 **Example: Unifying Knowledge Graph Embedding Models**

$$\begin{aligned} \text{HolE:} & \quad h \star r \quad (\text{circular correlation}) \\ \text{ComplEx:} & \quad \text{complex-valued embeddings} \\ \text{DistMult:} & \quad \text{real-valued symmetric embeddings} \end{aligned}$$

These can be understood as special cases of a more general framework.

ANALOGY Framework:

The ANALOGY framework unifies these approaches via analogical inference:

$$h + r \approx t$$

where an entity h combined with a relation r approximates the target entity t .

Interpretation: Different models correspond to different ways of representing this analogy.

Interpretability

Modern neural-KR systems (KG embedding models) work well, but it's unclear *why*.

Example 1.30 Example: Lack of Interpretability in Neural Models

Prediction: Neural model predicts:

Aspirin $\rightarrow$ treats $\rightarrow$ Headache

Question: Why? Which reasoning path?

Reasoning: Just learned embedding similarity.

There are different ways this could have been reasoned:

- **Path-based explanation:**

Aspirin $\rightarrow$ type $\rightarrow$ Painkiller $\rightarrow$ treats $\rightarrow$ Pain $\rightarrow$ symptom_of $\rightarrow$ Headache

- **Attention visualization:**

Model focuses on “anti-inflammatory” and “pain-relief” relations

- **Rule extraction:**

- IF drug $\rightarrow$ has_property $\rightarrow$ analgesic AND disease $\rightarrow$ has_symptom $\rightarrow$ pain
- THEN drug $\rightarrow$ treats $\rightarrow$ disease

And different applications:

- **Medical diagnosis:**

Doctor needs to see reasoning path before trusting AI recommendation

- **Legal applications:**

Must explain why system reached a conclusion

Scalability

Knowledge graphs can be extremely large which leads to a trade-off between and computational efficiency and model expressiveness.

Limited amount of works were applied to/tested on more than 1 million entities and inference over probabilistic logic is computationally intensive.

Knowledge Aggregation

Knowledge aggregation is the core of knowledge-aware applications. Many approaches rely on neural architectures. NLP has been significantly improved by large-scale pre-training. A key challenge remains the efficient and interpretable aggregation of knowledge.

Example 1.31 Query: Who directed Inception?

Knowledge neighbors:

[Nolan, Warner Bros, 2010, Sci-Fi, DiCaprio]

Attention weights:

director = 0.8, actor = 0.15, studio = 0.05

Interpretation: The model focuses on the most relevant relation (director).

Automatic Construction and Dynamics

Knowledge graphs are currently mostly manually constructed or extracted from structured sources (e.g. Wikipedia), which is very labor-intensive.

We would need automatic construction from large-scale unstructured content, but there are several challenges:

- Multimodality: Information can be in text, images, video
- Heterogeneity: Different sources use different schemas
- Large-scale: Billions of documents

Additionally, **knowledge graphs** should be dynamic, as real-world facts change over time.

Example 1.32 Static representation:

(Pluto, classified_as, Planet)

Dynamic representation:

(Pluto, classified_as, Planet, [1930-2006])

(Pluto, classified_as, Dwarf_Planet, [2006-present])

Requirement: System must track temporal validity of facts.

1.6 | Conclusion brought to you by *claude.ai*

“Knowledge representation remains a central discipline in artificial intelligence. While the field has historically undergone several paradigm shifts (from pure symbolic AI to knowledge graphs to hybrid systems), the fundamental motivation remains unchanged: How can we structure complex knowledge in a form that machines can understand

and use?

The answer continues to evolve. Modern systems combine:”

- Expressiveness of logical systems
- Efficiency of algorithmic approaches
- Scalability of big data techniques
- Learning capability of neural networks
- Interpretability through explicit knowledge

1.7 | Course structure

foundation of KR

Knowledge graphs

KR in mL

grading:

50% of the exercise points (full portfolio) (Abgabe in 3er Gruppen(?)) present one exercise in tut present one paper in seminar 40–45 min (3er gruppen) final report 5–10 pages (50% marks on final report)

Es wird ein Moodle geben

1.8 | Formalities

CHAPTER 2

Logical Foundations: First-Order Logic (FOL)

Knowledge representation requires a formal system that expresses knowledge precisely and derives consequences from that knowledge.^[2]

The most relevant formal calculations are **propositional Logic** and **first-order logic**.

Definition 2.1 **propositional Logic**

todo

Example 2.1 “If a person is rich, they have a nice car”

That statement becomes

$$\begin{aligned} BobIsRich &\Rightarrow BobHasNiceCar \\ JaneIsRich &\Rightarrow JaneHasNiceCar \\ &\vdots \end{aligned}$$

Having to define every rule for every person is impractical and unintuitive though.

2.1 | Syntax

6 - Objects sind x,y, ... 8 - not a well-formed formula weil es keinen Sinn macht

14 - unary predicate, welche gibt es noch??

Definition 2.2 **formula**

todo

Definition 2.3 **atomic formula**

todo

Definition 2.4 **compound formula**

todo

Definition 2.5 **quantified formula**

todo

Example 2.2 **Building a Formula – Example from OLP**

Definition 2.6 **free**

todo

Definition 2.7 **bound**

todo

Example 2.3 Von f 18

Definition 2.8 **sentence**

todo

Example 2.4 Von f 19

Definition 2.9 **substitution**

todo

Example 2.5 Von f 20.1

Definition 2.10 **captured**

todo

Example 2.6 Von f 21

2.2 | Semantics

Definition 2.11 **language structure**

todo

Definition 2.12 **domain**

todo

Definition 2.13 **valid**

todo

Definition 2.14 **satisfiable**

todo

Definition 2.15 **unsatisfiable**

todo

2.3 | Inference

Definition 2.16 **inference**

todo

Definition 2.17 **Modus Ponens**

todo

Definition 2.18 **Modus Tollens**

todo

Definition 2.19 **Universal Instantiation**

todo

Definition 2.20 **Existential Instantiation**

todo

Definition 2.21 **Prenex Normal Form**

todo

F 38 in step 1: Rule 1: $A \rightarrow B = \neg A \cup B$ Rule 2: $\neg\forall x(P(x)) \equiv \dots$ Ist es hier neg oder sim ??**Definition 2.22** **clause**

todo

Definition 2.23 **literal**

todo

Definition 2.24 **Conjunctive Normal Form**

todo

Definition 2.25 **Resolution-based prover**

todo

Definition 2.26 **resolution**

todo

2.4 | Proof Systems & Completeness

Definition 2.27 **Natural Deduction**

todo

Definition 2.28 **Sequent Calculus**

todo

Definition 2.29 **derivability relation**

todo

CHAPTER 3

Logic Programming and Automated Reasoning

CHAPTER 4

Structured Knowledge Representation

Part II

Knowledge Graphs

CHAPTER 5

Start into semantic networks and knowledge graphs reasoning

CHAPTER 6

The Resource Description Framework (RDF)

CHAPTER 7

Knowledge Graph Reasoning

CHAPTER 8

Knowledge Acquisition from Data

CHAPTER 9

Representation Learning

Part III

Knowledge Representation in Machine Learning

CHAPTER 10

Knowledge Graph Embeddings

CHAPTER 11

Graph Neural Networks for Knowledge Representation

Part IV

Applications

CHAPTER 12

Knowledge Tracing

CHAPTER 13

Knowledge in Large Language Models

CHAPTER 14

KG- Aware Applications

Backmatter

List of Figures

1.1	This is Tweety. What can we infer about Tweety with the information given below? (image from https://www.pexels.com/photo/pink-cockatoo-perched-in-natural-habitat-29732587/)	10
1.2	Semantic Network	11
1.3	This is Dr. Dokter. She wants to diagnose her patient. (image from https://www.pexels.com/photo/female-doctor-in-scrub-suit-7659874/)	11

Definitions

abductive inference Reasoning to arrive at the best explanation. From observations to hypothetical causes. 13

atomic formula todo. 23

bound todo. 24

captured todo. 24

clause todo. 25

compound formula todo. 23

Conjunctive Normal Form todo. 25

data todo. 8

declarative representation Declarative representations describe facts about the world as logical statements. 9, 10

derivability relation todo. 25

domain todo. 24

Existential Instantiation todo. 25

explicit knowledge Knowledge that is formally encoded in symbols, rules or graphs and can be processed by machines. 8

first-order logic todo. 9, 23

formula todo. 23

frame Concept designed to provide a structured way to capture the essential aspects of a situation, facilitating easier retrieval and manipulation of information. They are akin to schemas or blueprints that organize knowledge into manageable chunks. Attributes of the represented concept are represented by **slots** which have facets and default values. 12

frame-based representation todo. 10

free todo. 23

inductive inference Generalising examples into patterns or rules. 13

inference todo. 24

information todo. 8

knowledge todo. 8

Knowledge Acquisition Processes The process of incorporating knowledge into the knowledge base through strategies such as Interactive sessions with experts or Automatic generation from experience. 14

knowledge base todo. vi, 9

knowledge graph Graph-based knowledge representation where the entities are nodes and the relations are edges. Facts (subject, a predicate, and an object) are represented as triples.. 15, 16, 20

Knowledge Representation KR is a central problem in Artificial Intelligence that focuses on developing sufficiently precise notations for representing knowledge. 7, 17

language structure todo. 24

literal todo. 25

logical inference Inference that uses formal deduction from known facts. 12

logical representation scheme Representation schema based on formal logic and represents knowledge through logical formulas. It is strongly based on semantics and inference. 10

Modus Ponens todo. 24

Modus Tollens todo. 24

multi-hop todo. 18

Natural Deduction todo. 25

Natural Language Processing todo. 15

one-hop todo. 17

ontology todo. 8

preferred inference Reasoning based on typical assumptions, subject to certain reservations. 13

Prenex Normal Form todo. 25

procedural representation Representation schema that encodes knowledge as procedures or programmes that describe how to use that knowledge. It is efficient but hard to interpret. 10, 11

propositional Logic todo. 23

quantified formula todo. 23

representation scheme A representation scheme is a formal notation used to define a **knowledge base** (Mylopoulos 1980). A knowledge base contains facts and rules and serves as a model of the real world. A representation scheme is therefore the ‘language’ in which we express knowledge, much like programming languages define the syntax and semantics in which we write instructions.. 9

resolution todo. 25

Resolution-based prover todo. 25

satisfiable todo. 24

semantic network todo. 9

semantic representation scheme Graph-based representations in which nodes represent entities and edges represent relationships between them. It supports inheritance and classification. 11

Semantic Search todo. 16

sentence todo. 24

Sequent Calculus todo. 25

slot todo. v, 12

specialized procedures Fixed inference patterns, which are often implemented as routines. 12

substitution todo. 24

test todo. 15

Universal Instantiation todo. 25

unsatisfiable todo. 24

valid todo. 24

Bibliography

- [1] Jiahang Cao et al. “Knowledge Graph Embedding: A Survey from the Perspective of Representation Spaces”. In: *ACM Comput. Surv.* 56.6 (Mar. 2024). ISSN: 0360-0300. DOI: [10.1145/3643806](https://doi.org/10.1145/3643806). URL: <https://doi.org/10.1145/3643806>.
- [2] Randall Davis, Howard Shrobe, and Peter Szolovits. “What Is a Knowledge Representation?” In: *AI Magazine* 14.1 (1993), pp. 17–33. DOI: <https://doi.org/10.1609/aimag.v14i1.1029>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1609/aimag.v14i1.1029>. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1609/aimag.v14i1.1029>.
- [3] Aidan Hogan et al. “Knowledge Graphs”. In: *ACM Comput. Surv.* 54.4 (July 2021). ISSN: 0360-0300. DOI: [10.1145/3447772](https://doi.org/10.1145/3447772). URL: <https://doi.org/10.1145/3447772>.
- [4] *Knowledge Representation 01 - Knowledge Basic Introduction, Motivation, History*. Vol. 1. Lecture Notes Knowledge Representation.